

AI TRAINING WORKBOOK

Full Professional Version — Modules, Exercises, and Facilitator Guide

Document Type: Self-paced + facilitator-led training program

Audience: Staff training participants and facilitators

Version: 2.0 · Last Updated: May 2026

Estimated Read Time: Self-paced (8-12 hours total program)

Table of Contents

1. How to Use This Workbook
2. Module 1 — AI Fundamentals
3. Module 2 — Prompting Basics
4. Module 3 — Working with AI Safely
5. Module 4 — AI for Communication & Writing
6. Module 5 — AI for Research & Analysis
7. Module 6 — AI for Process Automation
8. Module 7 — Reviewing AI Output
9. Module 8 — Building Your AI Habit
10. Facilitator Guide
11. Completion Certificate Template

1. How to Use This Workbook

This workbook supports a structured 8-module AI training program. Each module is designed for 45-60 minutes of self-paced study followed by a short hands-on exercise. The facilitator guide at the end offers prompts and discussion points for groups working through the material together.

1.1 Learning paths

- **Self-paced individual:** work through one module per week. Total: 8 weeks.
- **Facilitator-led group (compressed):** two-day workshop, 4 modules per day.
- **Facilitator-led group (extended):** 8 weekly 90-minute sessions, module per session.
- **Coaching model:** one module as homework, one as group coaching call. Repeat for 8 weeks.

1.2 What you'll need

- Access to an approved enterprise AI tool (e.g., ChatGPT Team, Copilot, or equivalent).
- About 60 minutes per module, plus 30 minutes for the exercise.
- A notebook (paper or digital) for prompts and reflections.
- A sample of your actual work to use in exercises — emails, draft documents, reports.

1.3 How each module is structured

- **Setup:** The skill or concept and why it matters (5 minutes).
- **Core content:** Frameworks, techniques, and examples (20-30 minutes).
- **Demonstration:** Walked-through example using realistic work (10 minutes).
- **Practice:** Hands-on exercise using your own work (15-20 minutes).
- **Reflection:** Short questions to consolidate learning (5 minutes).
- **Assessment:** Quick check that you've absorbed the material.

2. Module 1 — AI Fundamentals

Learning objectives

- Understand what an LLM is and how it differs from traditional software.
- Articulate three things AI does well and three it doesn't.
- Recognize common terms: model, prompt, context window, hallucination, fine-tuning.
- Set expectations for your own AI use.

What AI actually is

At the simplest level, modern AI (specifically generative AI and LLMs) is software trained on enormous quantities of text to predict the next word given everything that came before. That sounds reductive, and it is, but the implications are profound: a system that learns to predict text well enough must internalize a lot about language, the world, and reasoning patterns.

What it's NOT: it doesn't 'understand' in a human sense, it doesn't have lived experience, it doesn't access the internet in real-time unless explicitly wired to, and it doesn't store conversations across separate sessions by default.

What AI does well

- **Drafting at scale.** Producing first drafts of writing, code, analyses faster than most humans.
- **Summarizing.** Condensing long content into shorter forms.
- **Pattern recognition in text.** Finding themes, extracting structured info from unstructured text.
- **Format transformation.** Email to outline, outline to slides, slides to talking points.
- **Brainstorming.** Generating options for you to pick from.
- **Language tasks.** Translation, tone adjustment, simplification.

What AI doesn't do well

- **Facts requiring recency.** Models have a training cutoff. They don't know what happened last week.
- **Math reliably.** They can do simple arithmetic but make confident errors at scale.
- **Counting and exact lookups.** Don't trust an LLM to count things in a document.
- **Citation.** They will make up plausible-sounding sources.
- **Privileged or proprietary knowledge.** If it wasn't in training data, it's guessing.
- **Long-term context.** A long conversation thread loses earlier context.

Glossary

Term	Meaning
LLM	Large Language Model — the underlying AI system
Prompt	Your input to the model

Context window	How much text the model can consider at once
Hallucination	Confident output that is factually wrong
Fine-tuning	Specializing a model on additional data
RAG	Retrieval-Augmented Generation — model with access to your documents
Agent	AI that can take actions, not just produce text
Token	Roughly a word fragment; usage is measured in tokens
Temperature	Setting that controls randomness in outputs
System prompt	Background instructions that shape how the AI behaves

Exercise 1

Try this

Open your approved AI tool. Ask it: 'Explain in plain English what you are and how you work. Then list five things you do well and five things you don't.' Note where its self-description matches what you read above, and where it differs.

Module 1 assessment

- Can you name three things AI does well? ✓
- Can you name three things AI doesn't do well? ✓
- Can you explain a 'hallucination' in one sentence to a colleague? ✓

3. Module 2 — Prompting Basics

Learning objectives

- Write prompts that produce useful output the first time.
- Use the C.R.E.A.T.E. structure for any non-trivial prompt.
- Recognize and fix five common prompt failures.

The C.R.E.A.T.E. structure

- **C — Context:** Who you are, what you're doing, what the audience cares about.
- **R — Role:** What expertise the AI should adopt.
- **E — Example:** Show what good output looks like.
- **A — Action:** Be specific about the task.
- **T — Tone:** Formal, casual, technical, warm, dry.
- **E — Format:** Bullet list, paragraph, table, JSON, etc.

Bad prompt vs. better prompt

Bad

'Write something about AI for my newsletter.'

Better

'You are a newsletter writer for an audience of small-business owners in Maine. Write a 300-word newsletter intro on a recent AI development. Tone: warm, plain-language, slightly skeptical. Format: 3 short paragraphs. Goal: get readers to click through to the full article. End with a 1-sentence question.'

Five common prompt failures

- 1 **Too short.** A 5-word prompt gives a generic answer.
- 2 **No audience.** Without specifying who reads the output, the AI defaults to vague middle ground.
- 3 **No example.** One example of what 'good' looks like is worth a thousand words of instruction.
- 4 **Buried instructions.** If your most important instruction is in the middle, the AI may miss it.
- 5 **Ambiguous success criteria.** 'Make it better' is meaningless. 'Cut 20% of the length, keep all numbers, simplify jargon' is actionable.

Iterating on a prompt

You usually won't get the perfect prompt on the first try. The skill is iterating. After the first response, ask: what's missing? What's wrong? What should be different? Then update the prompt with the additional instruction. Three iterations is usually enough.

Exercise 2

Try this

Pick a real task from your inbox or to-do list. Write a one-line prompt. Try it. Now rewrite using C.R.E.A.T.E. Try again. Compare the two outputs.

4. Module 3 — Working with AI Safely

The three rules

- **Rule 1: Classify before you paste.** Is this data public, internal, confidential, or restricted? The classification tells you which AI tool, if any, is appropriate.
- **Rule 2: Verify outputs that matter.** AI is a draft. You are the editor.
- **Rule 3: Disclose where required.** If your role or your audience expects to know AI was involved, say so.

Never share with public AI tools

- Customer or client identifying information (PII).
- Protected health information (PHI).
- Financial records, account numbers, payment details.
- Source code or proprietary algorithms.
- Confidential strategy or financial documents.
- Employee personal information.
- Vendor contracts.

What to do if you accidentally shared something

- 1 Stop. Don't try to delete by pasting more stuff to flush it.
- 2 Note exactly what you shared and with which tool.
- 3 Tell your supervisor and security team immediately. Faster is always better.
- 4 Document what data, what tool, what time. They will need this.
- 5 Follow internal incident response steps.

Important

Honest, immediate reporting is treated favorably. Hiding mistakes is treated as a separate violation. The org wants to know.

Phishing in the age of AI

- Grammatical perfection is no longer a 'this might be real' signal.
- Personalization in cold emails is now trivial to fake.
- Voice cloning makes 'I heard your boss say' attacks viable.
- Out-of-band verification is now mandatory for anything financial.

Exercise 3

Try this

Look at the last 10 emails you sent. For each, identify: would it be safe to paste this into a public AI tool? An approved enterprise AI? What classification does each one carry?

5. Module 4 — AI for Communication & Writing

Where AI shines for writing

- First drafts of emails, letters, memos.
- Tone-shifting (formal → casual, technical → plain).
- Length adjustment (expand a bullet into a paragraph, condense a page into bullets).
- Rewriting for a different audience.
- Generating multiple options to pick from.
- Proofreading and grammar.

Where it falls down

- Anything requiring your specific factual knowledge — meeting outcomes, customer specifics.
- Tone that requires reading the room.
- Difficult news that requires careful sequencing.
- Anything where authenticity matters more than polish.

Three useful patterns

Pattern A: Draft + edit

You give the AI bullet-point notes; AI produces a first draft; you edit. This is faster than starting from a blank page for most non-trivial writing.

Pattern B: Multiple options

Ask the AI for three different versions of the same message — different tones, different lengths, or different angles. You choose, sometimes mixing elements.

Pattern C: Critique

You write the draft. Then you ask the AI to critique it: what's unclear, what's redundant, what's missing, what's the most important thing in here. Then you edit based on the critique. Better outcomes than asking AI to rewrite.

Exercise 4

Try this

Take a difficult email you've been putting off. Use Pattern A: write five bullet points of what you want to say, then ask the AI to draft a 150-word email. Edit. Compare your time to write versus how long this would have taken from scratch.

6. Module 5 — AI for Research & Analysis

What works

- Brainstorming research questions.
- Summarizing a document you've uploaded.
- Extracting structured data from messy text (with verification).
- Comparing two or more documents.
- Drafting interview guides.
- Generating hypotheses to test.

What's risky

- **Citations.** Without RAG or a connected search tool, models will hallucinate plausible citations.
- **Statistics.** AI may invent or misquote numbers. Always verify against source.
- **Recent events.** Training cutoffs limit recency knowledge.
- **Specialized knowledge.** If it wasn't well-represented in training, you'll get confident wrong answers.

Verification practices

- 1 Spot-check every cited statistic against an original source.
- 2 If a source is named, find the actual link.
- 3 For specialized claims, consult a human expert before relying.
- 4 If you can't verify, treat as 'hypothesis to investigate' not 'fact to share'.

Exercise 5

Try this

Pick a topic you'd research for work. Ask the AI for five surprising facts about it, with sources. Now verify each one. Note how many turn out to be wrong, partially right, or unverifiable.

7. Module 6 — AI for Process Automation

Where automation makes sense

- Repetitive text generation: standard responses, follow-ups, summaries.
- Data extraction from forms, emails, receipts.
- Triage: classifying incoming requests into categories.
- Translation between formats: meeting notes → action items.

Where it doesn't

- When the process needs judgment that's hard to articulate.
- When errors are very costly.
- When the inputs are highly variable in structure.
- When the upstream process is broken (don't automate broken processes).

The automation worksheet

- What's the process? Describe in 3 sentences.
- How often does it run?
- What's the cost of a wrong outcome?
- What does a human currently catch that AI might miss?
- How will you measure if it's working?
- Who's accountable if it fails?

Exercise 6

Try this

Identify one task you do every week that takes 15+ minutes. Run it through the worksheet. Decide: is it a good automation candidate? If yes, pilot it for two weeks before committing.

8. Module 7 — Reviewing AI Output

The 5-question review

- 1 **Accuracy:** Is every fact in here true and verified?
- 2 **Completeness:** Is anything important missing?
- 3 **Tone:** Does this sound like me / my organization?
- 4 **Risk:** Could this cause harm or embarrassment if wrong?
- 5 **Attribution:** Do I need to disclose that AI was involved?

Common AI output failures to watch for

- Confident statements with no factual basis.
- Sentences that sound right but mean nothing.
- Stereotypes or assumptions about groups.
- Subtly wrong numbers presented confidently.
- Padding and repetition.
- Overly enthusiastic tone.
- Missing the actual question you asked.

Exercise 7

Try this

Take an AI-generated document you produced recently. Run the 5-question review on it. Note what you find.

9. Module 8 — Building Your AI Habit

The transition challenge

Most people don't fail to learn AI; they fail to integrate it. The hard part isn't learning to prompt — it's remembering to use AI at the moment a task lands on you.

Three integration habits

- **The 30-second check.** Before starting any text-heavy task, ask yourself: would AI help? If yes, start there.
- **The Friday review.** Once a week, list five tasks that took longer than they should have. Could AI have helped? Note the prompt to use next time.
- **The shared prompt library.** When you find a prompt that works well, save it. Share with colleagues.

Common adoption traps

- **Perfectionism.** Spending an hour crafting the perfect prompt to save 20 minutes of work.
- **Trust collapse.** One bad AI experience leading you to stop using it for everything.
- **Tool sprawl.** Trying every new AI tool and mastering none.
- **Hidden use.** Using AI but not telling anyone, so others can't learn from your wins.

Your 30-day plan

- Week 1: Use AI for at least one task every day. Note what works and what doesn't.
- Week 2: Save your three best prompts.
- Week 3: Find one task to automate or speed up significantly.
- Week 4: Share what you've learned with at least one colleague.

Exercise 8

Final exercise

Write down three commitments for what you'll do differently in the next 30 days. Share them with your manager or learning partner. Schedule a check-in in 30 days.

10. Facilitator Guide

10.1 Running a group workshop

- Have everyone log into the approved AI tool before starting.
- Demonstrate prompts on a shared screen for the first 2-3 modules.
- Use real workplace examples whenever possible.
- Allow time for everyone to share what worked and what didn't.
- Watch for participants who are hesitant or stuck; pair them with confident peers.

10.2 Discussion prompts per module

- **Module 1:** What surprises you about how AI actually works? What expectations have you had to update?
- **Module 2:** What's a prompt you've used that worked unexpectedly well?
- **Module 3:** What are the data categories in YOUR work? Where are the lines?
- **Module 4:** Which writing tasks do you most want to offload?
- **Module 5:** Has AI given you confidently-wrong information before? What happened?
- **Module 6:** What's a process in your team that's begging to be automated?
- **Module 7:** What's your review process going to look like?
- **Module 8:** What's the one change you're making starting tomorrow?

10.3 Common facilitator challenges

- **The skeptic:** hear their concerns; don't try to argue. Pair them with a power user.
- **The enthusiast:** direct their energy toward helping others; watch for over-confidence.
- **The disengaged:** ask what tasks they find most tedious; that's the entry point.
- **The anxious:** emphasize the policy and safety nets; they want to be safe more than fast.

11. Completion Certificate Template

Certificate of Completion

This certifies that **[Participant Name]** has completed the AI Impact Maine 8-module AI Training Workbook program, including:

- AI Fundamentals
- Prompting Basics
- Working with AI Safely
- AI for Communication & Writing
- AI for Research & Analysis
- AI for Process Automation
- Reviewing AI Output
- Building Your AI Habit

Awarded: _____

Facilitator: _____

About AI Impact Maine

AI Impact Maine helps Maine businesses, municipalities, and nonprofits assess, plan, and responsibly implement AI. We offer readiness assessments, training programs, policy and governance design, vendor evaluation, and ongoing strategic support.

Get in touch

Email: ahmad@aiimpactmaine.com

Web: aiimpactmaine.com

Book a free 30-minute consultation: aiimpactmaine.com/contact