

CHATGPT SAFETY GUIDE

Full Professional Version — Modules, Scenarios, and Reference Card

Document Type: Staff training + reference

Audience: All employees who use or might use ChatGPT or similar tools

Version: 2.0 · Last Updated: May 2026

Estimated Read Time: 75 minutes

Table of Contents

1. Why This Guide Exists
2. Module 1 — What ChatGPT Is
3. Module 2 — What You Can and Can't Share
4. Module 3 — Verifying What ChatGPT Tells You
5. Module 4 — Recognizing AI-Powered Scams
6. Module 5 — Privacy & Account Hygiene
7. Module 6 — Special Cases (Healthcare, Education, Legal)
8. Module 7 — When to Disclose AI Use
9. Module 8 — What to Do If Something Goes Wrong
10. Practical Scenarios
11. Facilitator Guide
12. Quick Reference Card

1. Why This Guide Exists

ChatGPT and similar tools (Claude, Gemini, Copilot) are now used by hundreds of millions of people. Many use them at work, often without explicit permission and without knowing what's risky. This guide is for the worker, not the IT department. It will not slow you down. It will keep you out of avoidable trouble.

Each module is 10-15 minutes. You can read straight through (about 75 minutes total) or take one a week.

What this guide will and won't tell you

- . **Will:** The most common ways people get into trouble with ChatGPT at work.
- . **Will:** How to spot AI-related scams aimed at you or your organization.
- . **Will:** What to do when you make a mistake.
- . **Won't:** Tell you not to use ChatGPT. It's a useful tool.
- . **Won't:** Replace your organization's specific policy. Your policy wins when there's a conflict.

2. Module 1 — What ChatGPT Is

What it actually does

ChatGPT generates text by predicting what should come next. It was trained on huge amounts of text from the internet and licensed sources. It doesn't 'know' things the way a person knows; it has patterns from training that often produce accurate text, but it's not a search engine and it's not a calculator.

Three common misconceptions

- **Misconception:** 'It knows everything on the internet.' Reality: it knows what was in training. Often a year or more old.
- **Misconception:** 'My conversations are private to me.' Reality: depends on the account type. Free accounts may be used for training.
- **Misconception:** 'If it says it with confidence, it's right.' Reality: confidence and accuracy are unrelated in AI output.

Free vs. Plus vs. Team vs. Enterprise

Tier	Who pays	Data used for training?
Free	Nobody (ad-free model)	Yes, by default
Plus	Individual	Maybe (toggle in settings)
Team	Company	No (per OpenAI terms)
Enterprise	Company	No (per OpenAI terms, with admin controls)

Key takeaway

If your org pays for a Team or Enterprise account, your conversations are not used for training and have stronger privacy guarantees. If you're using a free or Plus account for work, treat every input as potentially exposed.

3. Module 2 — What You Can and Can't Share

The simple rule

Before you paste anything, ask: 'Would I be comfortable if this ended up in a Reddit post, a leaked dataset, or a regulator's inbox?' If the answer is no, don't paste.

Never share with consumer ChatGPT

- Names + contact info of customers, clients, students, patients.
- Health information of any kind connected to a person.
- Account numbers, payment details.
- Internal financial reports.
- HR records: discipline, compensation, evaluations.
- Personnel issues involving real people.
- Source code of internal products.
- Confidential strategy documents.
- Vendor contracts.
- Anything labeled 'confidential' or 'restricted' in any system.

Usually safe to share

- Public marketing copy you wrote.
- Generic templates you want polished.
- Industry trends without organization-specific data.
- Generic technical questions ('how does X work?')
- Your own learning questions.

De-identification doesn't make it safe

Replacing names with X doesn't always work. If the rest of the document is specific enough (dates, dollar amounts, locations, project names), re-identification is often trivial. Treat de-identification as a courtesy, not a safety mechanism.

What if my work IS confidential?

If your work requires confidentiality (legal, healthcare, sensitive client work, etc.), you should use an enterprise-grade tool with appropriate contractual protections — or not use AI at all for that work. Check your organization's approved tool list.

4. Module 3 — Verifying What ChatGPT Tells You

The hallucination problem

ChatGPT will sometimes produce confident, plausible, completely wrong answers. This is not a bug — it's how the underlying technology works. Your job is to verify before relying.

Things ChatGPT often gets wrong

- **Citations.** Made-up papers, court cases, books.
- **Specific statistics.** Numbers that look right but aren't.
- **Recent events.** Training cutoff means recent news is unreliable.
- **Specialized fields.** Niche topics often elicit hallucinations.
- **Names + dates + places.** Conflated or invented.
- **Counting and arithmetic at scale.** Confident but often wrong.

The 30-second verification habit

- 1 If a fact will be sent externally or relied on for a decision, verify.
- 2 Spot-check by Googling the specific claim.
- 3 If a citation is given, find the actual source.
- 4 If you can't find a source, treat the claim as unverified.
- 5 For numbers, find the primary source — government data, official reports.

When 'good enough' is good enough

- Internal brainstorming notes.
- First drafts you'll edit heavily.
- Learning explorations for your own understanding.
- Personal use you'll fact-check before relying on.

5. Module 4 — Recognizing AI-Powered Scams

AI has changed the scam economy. Tactics that used to be amateurish are now polished. This module covers the patterns to watch for.

5.1 Phishing emails

- Perfect grammar is no longer reassuring. Scammers use AI to write.
- Personalization is easy: scammers can scrape LinkedIn + your company site and produce a tailored email cheaply.
- Watch for: urgency, unusual sender domains, unexpected attachments, requests to change payment info.
- Verify any financial or sensitive request through a SECOND channel (phone call to a known number, in-person, etc.).

5.2 Voice clone calls

AI can clone a voice from ~30 seconds of audio. Don't trust voice alone for high-value transactions.

- Establish a code word with family for emergency calls.
- Establish a verification step with finance/HR for transfers.
- If a call feels off (rushed, threatening, asking for secrets), hang up and call back on a known number.

5.3 Deepfake video calls

Video deepfakes are good enough now to fool casual verification. The Hong Kong \$25M case (an employee was tricked by a fake CFO on a video call) is the canonical example.

- If anything feels off, pause. Real executives can wait 10 minutes.
- Schedule large transactions, don't react on-the-call.
- Multi-channel verification for big-money decisions.

5.4 Fake job offers / fake applicants

- AI-generated job offers targeting your employees: meant to harvest credentials or hijack accounts.
- AI-generated applicants applying to your roles: increasingly hard to spot.
- Live interview steps for anything beyond entry-level help.

6. Module 5 — Privacy & Account Hygiene

Your ChatGPT account

- Use a strong, unique password. Use a password manager.
- Enable two-factor authentication.
- Don't share your account with coworkers.
- Use a work email for work accounts, personal email for personal.
- Review privacy settings: 'Improve the model for everyone' should be OFF for work use.

Chat history

- Default behavior: chat history is saved.
- You can turn this off in settings.
- Even with history off, conversations may be retained for 30 days for abuse review.
- For most enterprise tiers, behavior differs — check with IT.

Custom GPTs and shared instructions

- Anything you put in 'custom instructions' or a custom GPT is part of every chat.
- Don't put proprietary info in custom instructions.
- Treat shared custom GPTs the same way you'd treat a shared inbox.

7. Module 6 — Special Cases

7.1 Healthcare

- Never paste PHI into consumer ChatGPT. This is a HIPAA violation.
- Even patient initials + condition is usually identifiable.
- Approved medical-AI tools have BAAs in place; consumer tools do not.
- When in doubt, ask your compliance officer.

7.2 Education

- FERPA protects student records. Don't paste student names or details.
- Different rules for K-12 vs. higher ed; check your institution's policy.
- Consider parental consent issues for K-12.
- Using AI to grade has additional fairness considerations.

7.3 Legal

- Privilege can be compromised by sharing with consumer AI.
- Multiple courts have sanctioned lawyers for citing AI-hallucinated cases.
- Use legal-specific tools with appropriate protections.
- Always verify any cited case or statute against an actual source.

7.4 Financial Services

- PCI-DSS applies to payment data.
- SEC rules apply to material non-public information.
- Customer financial details are restricted.
- Use approved enterprise tools only.

8. Module 7 — When to Disclose AI Use

When to definitely disclose

- When your role or contract requires it.
- When your audience would object to learning afterward that AI was involved.
- When you're publishing creative work.
- When AI is making decisions about people (employment, services, finances).
- When required by law or regulation.

When disclosure is optional

- When AI helped you draft something you then heavily edited.
- When AI did research that you verified independently.
- When AI helped you brainstorm ideas you then evaluated.
- Internal-use materials where everyone knows AI is part of the toolkit.

How to disclose

- Brief footer note: 'Drafted with AI assistance; reviewed and edited by [name].'
- In context: 'I asked AI to brainstorm some options. Here are the ones I selected.'
- Don't oversell or undersell. Be accurate.

9. Module 8 — What to Do If Something Goes Wrong

You shared confidential info with consumer AI

- 1 Stop using the tool for that data immediately.
- 2 Don't try to 'delete' by pasting more text. That doesn't work.
- 3 Note exactly what was shared, when, with which tool.
- 4 Tell your manager and your security/IT team within the hour.
- 5 Help them assess the impact and notification obligations.

You relied on a hallucinated fact

- 1 Confess promptly. Better to surface than have someone else find out.
- 2 Issue correction to whoever received the wrong info.
- 3 Update the source if it was a publication or report.
- 4 Learn the verification habit (Module 3).

You think you've been targeted by an AI scam

- 1 Don't act on the request. Don't transfer money, share credentials, or click links.
- 2 Verify through a known channel (call the person directly on a known number).
- 3 Report to IT/security with details.
- 4 Help warn others if it's a pattern affecting multiple people.

You see a coworker doing something risky

- Friendly heads-up is usually best first step.
- If they don't change behavior or it's a serious risk, escalate to your manager or IT.
- Don't shame publicly — undermines reporting culture.

10. Practical Scenarios

Scenario 1: The customer escalation

Setup

A customer sends you a long, angry email. You want to use ChatGPT to draft a calm, professional response. Is this OK?

Likely OK if (a) you redact the customer's name and any identifying details, OR (b) you use an approved enterprise AI. Even better: use AI to brainstorm responses you then write yourself. Avoid pasting the full email + signature + account info.

Scenario 2: The board memo

Setup

You need to write a memo for the board summarizing Q3 financial results. The data is in a spreadsheet you have access to. Is ChatGPT helpful here?

Be careful. Don't paste actual financials into consumer ChatGPT. You can: ask it for memo structure suggestions, ask it to clean up your prose after you write it, or use an approved enterprise AI with appropriate data handling. Never paste pre-release financial data.

Scenario 3: The unsolicited email

Setup

You get an email from someone claiming to be a vendor, with a contract attached, asking for a wire transfer to a new account. The email is well-written and references a real project.

This is the modern phishing pattern. Stop. Don't send the money. Call the vendor on a known phone number (not from the email). Verify out-of-band. Forward suspicious email to your security team.

Scenario 4: The job applicant

Setup

You receive a polished, perfectly-targeted resume and cover letter from someone applying for an open role. Their LinkedIn profile is sparse.

Could be legit, could be AI-generated. Include a live conversation step in your screening. Ask questions that require specifics about claimed experience. Be especially careful with remote roles where you never meet the person.

Scenario 5: The 'helpful' chrome extension

Setup

A colleague recommends a Chrome extension that 'auto-summarizes' every web page or email you open using AI.

Don't install without IT review. The extension is reading everything you read. Where does that data go? Who else has access? Stick to approved tools.

11. Facilitator Guide

Running a 2-hour session

- Skim modules 1-2 (15 min).
- Work through modules 3-5 in depth (60 min).
- Group discussion on scenarios (30 min).
- Review the reference card (15 min).

Running an 8-week schedule

- One module per week, 15-30 minutes.
- Lightweight discussion at start of each session.
- Final session: review scenarios and reference card.

Common questions you'll hear

- **'Is this banned?'** Usually no. Discourage rule-as-ban framing. The point is sensible use.
- **'How do I know what's confidential?'** Refer to your data classification scheme. If you don't have one, default to the categories in this guide.
- **'But everyone's using ChatGPT.'** Yes. The question is whether they're using it well. This training helps them do that.

12. Quick Reference Card

PRINT THIS PAGE

Pin to your monitor or keep in your notebook.

Never paste into consumer ChatGPT:

- Customer / client / patient / student names + details
- Health information
- Financial records, account numbers
- Proprietary code or strategy docs
- Personnel information
- Vendor contracts

Always do before relying on AI output:

- Verify the facts (especially numbers and citations)
- Check the tone is right
- Edit; don't paste unmodified
- Disclose if your role/audience requires it

Watch out for AI scams:

- Perfect-grammar phishing
- Voice clone urgency calls
- Deepfake video meetings
- Polished but suspicious job applicants

If something goes wrong:

- Report fast — speed matters more than perfection
- Tell exactly what happened
- Honest reporting protects you and your colleagues

About AI Impact Maine

AI Impact Maine helps Maine businesses, municipalities, and nonprofits assess, plan, and responsibly implement AI. We offer readiness assessments, training programs, policy and governance design, vendor evaluation, and ongoing strategic support.

Get in touch

Email: ahmad@aiimpactmaine.com

Web: aiimpactmaine.com

Book a free 30-minute consultation: aiimpactmaine.com/contact