

PROMPT ENGINEERING GUIDE

Full Professional Version — Techniques, Templates, and Patterns

Document Type: Practitioner's reference for effective AI prompting

Audience: Anyone using LLMs as part of their work

Version: 2.0 · Last Updated: May 2026

Estimated Read Time: 60 minutes (skim) or 3 hours (work through every section)

Table of Contents

1. Why Prompting Matters (And Where It Doesn't)
2. The 10 Core Techniques
3. Advanced Patterns Library
4. Troubleshooting Guide
5. 50+ Prompt Templates
6. Department-Specific Collections
7. Building Your Personal Prompt Library
8. Common Mistakes and How to Fix Them
9. When to Stop Prompting and Start a New Chat

1. Why Prompting Matters (And Where It Doesn't)

Prompt engineering is the practice of writing instructions that get useful, reliable output from AI models. The right prompt can turn a generic response into something genuinely useful for your specific situation. The wrong prompt produces output that sounds smart and isn't.

1.1 When prompting heavily matters

- When the output has a specific format or audience.
- When you need consistent quality across similar tasks.
- When you're building a reusable workflow.
- When accuracy in domain-specific facts is required.

1.2 When prompting matters less

- When you're brainstorming and quantity of options matters more than precision.
- When you're going to heavily edit the output anyway.
- When the task is trivial: 'summarize this paragraph in 3 sentences.'

1.3 What good prompting won't fix

- Lack of relevant training data on a topic.
- Math errors beyond simple arithmetic.
- Outdated knowledge past the training cutoff.
- The need for current real-world data the model can't access.

2. The 10 Core Techniques

2.1 Technique 1 — Role Assignment

Telling the AI who to be focuses its responses on the expertise you want.

Pattern

"You are a [role] with [years of experience] in [domain]. You are known for [specific quality]."

Example

"You are a senior HR director with 15 years of experience at small nonprofits. You are known for finding the humane phrasing in difficult conversations."

2.2 Technique 2 — Audience Specification

Telling the AI who's reading the output shifts vocabulary, tone, and structure.

Pattern

"This is for [audience]. They care about [priorities]. They have [knowledge level]."

2.3 Technique 3 — Output Format Specification

Specify the exact format you want. Don't assume the model will guess right.

- "Use bullet points."
- "Output as a table with columns X, Y, Z."
- "Three paragraphs of 100-150 words each."
- "Return JSON with these fields: title, summary, action_items."

2.4 Technique 4 — Examples (Few-Shot)

Show the model what good output looks like. One example is worth a paragraph of instruction. Two are usually enough.

Pattern

"Here's an example of the kind of output I want:

[example]

Now produce similar output for [new input]."

2.5 Technique 5 — Constraints

Tell the AI what NOT to do.

- "Don't use the word 'leverage.'"
- "Don't include any URLs or citations."
- "Don't exceed 150 words."
- "Don't editorialize. Just summarize."
- "Don't use any emojis."

2.6 Technique 6 — Step-by-Step Thinking

For complex tasks, ask the AI to think through it step by step before producing the final output.

Pattern

"Before producing the final answer, think through:

1. What information do I have?
2. What's the goal?
3. What are the key considerations?

Then give me your answer."

2.7 Technique 7 — Multiple Options

Ask for several variations. Pick or combine.

Pattern

"Give me three different versions:

1. Formal and direct
2. Warm and conversational
3. Brief and urgent"

2.8 Technique 8 — Critique Mode

After getting a draft, ask the AI to critique it.

Pattern

"Now critique this draft. What's unclear? What's redundant? What's the most important missing element?"

2.9 Technique 9 — Iterative Refinement

Treat prompting as a conversation, not a single-shot request.

- First pass: get something on the page.
- Second pass: tell it what's wrong.
- Third pass: tighten the specific weaknesses.
- Usually 3-4 iterations beats a single perfect prompt.

2.10 Technique 10 — Self-Verification

Ask the AI to check its own work for specific issues.

Pattern

"Now verify: (a) every fact is accurate, (b) every name is spelled right, (c) every number is consistent with the source, (d) nothing was made up."

Note: this isn't reliable on its own; it's a defense-in-depth measure. Still verify externally for anything that matters.

3. Advanced Patterns Library

3.1 Chain-of-thought

For reasoning-heavy tasks, ask the model to show its work.

Pattern

"Walk through your reasoning step by step, then give the final answer."

3.2 Self-consistency

For high-stakes answers, ask the model to solve the problem three different ways and check for consistency.

3.3 Persona conditioning

Specifying a real or composite persona produces more consistent stylistic output than abstract instructions.

Example

"Write in the voice of a small-town newspaper editor who's seen it all. Skeptical but not cynical."

3.4 Format-then-fill

First ask the model to design the structure (outline, table, schema). Then fill it in.

3.5 Constrained creativity

Hard constraints often produce more creative output than open-ended requests.

Example

"Write a 6-word product description." produces sharper output than "Write a short product description."

3.6 The 'translate to' pattern

Use AI to translate between forms: idea → outline → bullet → email → memo → speech.

3.7 The negative example pattern

Show the model what bad output looks like, so it knows what to avoid.

3.8 The role-rotation pattern

Ask the model to take multiple roles in sequence: 'Now critique the previous draft as a skeptical donor would.'

3.9 The audience-simulation pattern

Ask the model to react to your message as the audience would. Useful for anticipating questions.

3.10 The redact-and-share pattern

For confidential workflows, pre-process your input to remove identifying details, ask the model to work on the structure, then map results back to actual identifiers locally.

4. Troubleshooting Guide

4.1 Output is too generic

- Add audience specification.
- Add role assignment.
- Add concrete examples.
- Add specific constraints.

4.2 Output is too long

- Specify exact word count.
- Specify number of bullets or paragraphs.
- Ask for the 'sharpest 100 words.'

4.3 Output is too short or vague

- Ask for more specificity: 'Add concrete examples.'
- Ask for evidence: 'Support each claim with reasoning.'
- Request a target length.

4.4 Output is wrong tone

- Provide an example of correct tone.
- Use persona conditioning.
- Iteratively refine with specific tone notes.

4.5 Output is made-up / hallucinated

- Ask for sources and verify each.
- Add 'if you don't know, say so' instructions.
- Switch to a model with web access if available.
- Verify externally regardless.

4.6 Output ignores instructions

- Put critical instructions at the start AND end of your prompt.
- Use clear formatting: numbered lists, headers, capitalization for emphasis.
- Reduce instruction count — too many competing constraints confuse models.
- Break the task into stages with separate prompts.

4.7 Conversation has gone off the rails

- Start a new chat. Context windows degrade over long conversations.

- Bring the relevant context forward explicitly.
- Don't try to course-correct an exhausted conversation.

5. 50+ Prompt Templates

Direct, ready-to-use templates. Adjust the [bracketed] sections.

5.1 Email & Communication

- **Difficult email:** "Help me draft an email to [recipient role] about [topic]. Tone: professional, warm, not defensive. Goal: [specific outcome]. Length: under 200 words."
- **Follow-up:** "Draft a follow-up to this thread [paste thread]. Acknowledge their question. Confirm next step. Don't apologize for the delay."
- **Decline politely:** "Draft a decline of [opportunity]. Be respectful. Keep door open for future. Don't over-explain."
- **Request that doesn't sound demanding:** "Rewrite this request to sound collaborative, not demanding: [paste]."
- **Conflict de-escalation:** "This message [paste] is escalating a conflict. Rewrite my reply to lower the temperature without conceding the substance."
- **Acknowledgment of bad news:** "Draft a message acknowledging [bad event]. Be honest. Don't minimize. Indicate next steps."

5.2 Writing & Editing

- **Tighten:** "Cut this by 30% without losing any substance: [paste]."
- **Simplify:** "Rewrite for an 8th-grade reading level: [paste]."
- **Specify:** "Replace vague phrases with concrete details: [paste]."
- **Restructure:** "Reorganize this so the main point comes first: [paste]."
- **Tone shift:** "Rewrite from corporate to conversational: [paste]."
- **Critique:** "Critique this draft. What's unclear, redundant, missing? Don't rewrite — just critique."
- **Multiple openings:** "Give me five different opening lines for this piece."
- **Multiple closings:** "Give me three different closing paragraphs — one provocative, one reflective, one practical."

5.3 Analysis & Research

- **Document summary:** "Summarize this [doc] in: (1) one sentence, (2) three bullets, (3) one paragraph."
- **Compare:** "Compare these two [items]. Highlight three meaningful differences and two surprising similarities."
- **Synthesize:** "Across these documents [paste], identify common themes and disagreements."
- **Hypothesize:** "Given this data [paste], generate three plausible explanations. Mark each as well-supported, somewhat supported, or speculative."
- **Sanity-check:** "Look at this analysis [paste]. What's the weakest claim? Where would an expert push back?"

5.4 Meetings

- **Meeting summary:** "Summarize this transcript [paste] in: agenda, decisions, action items (with owners), open questions."
- **Pre-read for meeting:** "I'm meeting with [person] about [topic]. Suggest 5 questions I should ask and 3 things I should bring."
- **Post-meeting follow-up:** "Draft a follow-up email for this meeting [notes]. Include: thanks, summary of decisions, next steps, timeline."

5.5 Brainstorming

- **20 ideas:** "Give me 20 ideas for [problem]. Don't filter for quality — quantity first."
- **Adjacent angles:** "What's an angle on [topic] I'm probably missing?"
- **Devil's advocate:** "Argue against my plan [paste]. Find the three strongest objections."
- **Other industries:** "How would [other field] approach this problem?"

5.6 Decision support

- **Pros and cons:** "For [decision], list pros, cons, and unknowns. Then identify which factor matters most."
- **Reversibility check:** "Is this decision reversible? What would it cost to reverse if it goes wrong?"
- **Pre-mortem:** "Imagine this project failed in 6 months. What likely caused the failure? List the top 3."
- **Stakeholder map:** "For this decision [paste], who's affected, who decides, who needs to know?"

5.7 Learning

- **Explain like I'm 5:** "Explain [concept] as if I'm five years old. Use one analogy."
- **Explain technically:** "Explain [concept] to someone with strong general knowledge but not specialist training."
- **Compare and contrast:** "Compare [A] and [B]. Make the differences vivid."
- **Quiz me:** "Quiz me on [topic]. Five questions of varying difficulty. Wait for my answers before grading."

5.8 Code & Technical

- **Explain code:** "Walk me through this code [paste]. What does each section do?"
- **Refactor:** "Refactor this code for readability. Don't change functionality. Comment any non-obvious choices."
- **Debug help:** "This code [paste] is producing [error]. What are the three most likely causes? What should I check first?"
- **Write test cases:** "Write test cases for this function [paste]. Cover happy path, edge cases, and likely failures."

6. Department-Specific Collections

6.1 For executives

- Board memo prep, talking points generation.
- Pre-meeting briefings.
- Executive summary of long reports.
- Decision frameworks.
- Crisis communication drafts.

6.2 For HR

- Job description drafts.
- Interview question generators.
- Onboarding email sequences.
- Performance review feedback drafts (with caution).
- Policy explanation in plain language.

6.3 For development / fundraising

- Donor thank-you templates.
- Grant proposal first drafts.
- Campaign messaging variations.
- Donor segment-specific appeals.
- Major gift conversation prep.

6.4 For program staff

- Curriculum drafts.
- Client communication templates.
- Program report structure and drafts.
- FAQ generation from real inquiries.
- Resource translation.

6.5 For finance & operations

- Budget narrative drafts.
- Procurement memo generation.
- Vendor evaluation framework.
- Process documentation.
- Variance explanation drafts.

6.6 For marketing & communications

- Social media calendars.
- Blog post drafts.
- Press release structure.
- Newsletter content.
- Tagline brainstorming.

7. Building Your Personal Prompt Library

The hardest part isn't finding good prompts; it's finding them again when you need them. Build a library.

7.1 Structure

- Notes app, doc, or note-taking tool — whatever you'll actually use.
- Folders by use case: emails, analysis, meetings, etc.
- Each entry: prompt + context (when to use) + sample output.
- Tag the ones that have worked particularly well.

7.2 Cadence

- When you find a prompt that works, save it within the day.
- Friday review: anything from this week worth saving?
- Monthly: clean up; archive ones you don't use.
- Quarterly: share with colleagues. Steal theirs.

7.3 What makes a prompt worth saving

- It produces consistently useful output across similar tasks.
- It's longer or more nuanced than you'd reconstruct from memory.
- It encodes hard-won lessons (the right constraints, the right tone).
- Other people in your role could use it too.

8. Common Mistakes and How to Fix Them

- **Over-prompting:** too many instructions, model gets confused. Fix: prune to essentials.
- **Under-prompting:** too vague, output is generic. Fix: add audience + role + format.
- **Misplaced trust:** using AI output without verification. Fix: verify anything that matters.
- **Single-shot expectation:** hoping for perfection on the first try. Fix: iterate.
- **Forgetting context:** AI doesn't remember earlier chats. Fix: bring context with you or start fresh.
- **Apologizing:** 'please' and 'sorry' don't help. Just be clear and specific.
- **Fighting the model:** if it's getting a request wrong repeatedly, restructure the prompt or start fresh.

9. When to Stop Prompting and Start a New Chat

- The output is getting worse, not better, with each iteration.
- The model keeps repeating the same mistake despite correction.
- You've been at this for more than 15 minutes without progress.
- The model seems 'stuck' in a particular framing.
- You realize halfway through that the original prompt was the wrong shape.
- The chat is over 30 turns and important early context might be lost.

Starting a new chat is often faster than rescuing a bad one. Bring forward only the relevant context.

About AI Impact Maine

AI Impact Maine helps Maine businesses, municipalities, and nonprofits assess, plan, and responsibly implement AI. We offer readiness assessments, training programs, policy and governance design, vendor evaluation, and ongoing strategic support.

Get in touch

Email: ahmad@aiimpactmaine.com

Web: aiimpactmaine.com

Book a free 30-minute consultation: aiimpactmaine.com/contact