

RESPONSIBLE AI FRAMEWORK

Full Professional Version — Principles, Practices, and Playbooks

Document Type: Strategic framework for responsible AI adoption

Audience: Executive teams, board members, ethics committees, compliance leads

Version: 2.0 · Last Updated: May 2026

Estimated Read Time: 55 minutes

Table of Contents

1. What Responsible AI Means in Practice
2. The Six Principles
3. Governance Framework (15-page version)
4. Risk and Bias Assessment Toolkit
5. Stakeholder Engagement Playbook
6. Implementation Roadmap Template
7. Quarterly Review Templates
8. Maine-Specific Considerations
9. Self-Assessment Scorecard

1. What Responsible AI Means in Practice

Responsible AI is the deliberate practice of building, deploying, and operating AI systems in ways that are safe, fair, transparent, and accountable. The word 'practice' matters more than the word 'principles' — responsible AI is not a position you take, it's a set of habits an organization develops.

This framework distinguishes responsible AI from three things it is sometimes confused with: AI ethics (philosophy), AI safety (technical research on long-term risks), and AI compliance (meeting current regulations). All three intersect with responsible AI, but none is the same thing.

1.1 Why 'responsible' is doing work in this phrase

- **Responsibility implies accountability.** Someone is answerable for outcomes.
- **Responsibility implies foresight.** Anticipating downstream effects, not just reacting.
- **Responsibility implies trade-offs.** Choosing what to do AND what not to do.
- **Responsibility is organizational, not just individual.** Systems and incentives matter.

1.2 What this framework will NOT do

- Resolve open philosophical debates in AI ethics.
- Substitute for legal compliance review.
- Make responsible AI decisions easier — they are often genuinely hard.
- Promise specific business outcomes from responsible practice.

2. The Six Principles

This framework adopts six operational principles. They are deliberately chosen for usefulness, not novelty — they overlap with NIST AI RMF, OECD AI Principles, and most national frameworks.

2.1 Principle 1 — Human Oversight

Consequential AI decisions should have meaningful human review. 'Meaningful' is the key word: a rubber-stamp checkbox is not oversight.

- Define which decisions are 'consequential' for your context (affecting employment, housing, services, finances, safety).
- Specify what 'human review' means in practice: timeframes, qualifications, escalation paths.
- Audit: how often does the human reviewer actually change the AI output? If it's near zero, oversight may be cosmetic.

2.2 Principle 2 — Fairness

AI systems should not produce systematically worse outcomes for protected or vulnerable groups.

- Define 'fairness' for the context. Different definitions are sometimes mathematically incompatible.
- Pre-deployment: bias testing on historical data.
- Post-deployment: ongoing disparate-impact monitoring.
- Have a plan for what you'll do if disparate impact emerges.

2.3 Principle 3 — Transparency

People affected by AI should know it's there and roughly how it works.

- Disclose AI use to customers, employees, and the public as appropriate.
- Provide plain-language explanations of what the AI does and doesn't do.
- Maintain a public AI inventory if you serve the public.

2.4 Principle 4 — Privacy

Personal data deserves protection that doesn't degrade just because AI is involved.

- Apply existing privacy controls to AI workflows.
- Be especially careful about data flowing to AI training.
- Honor data subject rights (access, deletion) even when data is in AI systems.

2.5 Principle 5 — Security

AI systems are attack surfaces. Treat them accordingly.

- Apply standard security practices: encryption, access control, logging.
- Add AI-specific defenses: prompt injection protection, model integrity, output monitoring.
- Plan for incidents that are AI-specific.

2.6 Principle 6 — Accountability

Someone is responsible for AI outcomes in your organization. That someone is named and empowered.

- Single named accountability owner for the AI program.
- Operational owners for each deployed system.
- Documented decision trails so accountability can be exercised.
- Consequences for violations and rewards for surfacing issues.

3. Governance Framework

3.1 Governance versus management

Governance is direction-setting and oversight. Management is execution. A good governance framework lets management do its job without re-litigating every decision at the top.

3.2 The four governance functions

- **Direct:** Set strategy, risk appetite, and priorities.
- **Enable:** Provide resources, structure, and authority.
- **Oversee:** Monitor results against expectations.
- **Adjust:** Update direction as the environment changes.

3.3 Governance bodies

Body	Cadence	Function
Board / governing body	Quarterly	Strategic direction, risk appetite, escalation
Executive AI sponsor	Monthly	Resource allocation, decisions above committee authority
AI Review Committee	Monthly	Vendor approval, policy, incident review
Working groups	As needed	Specific projects, pilots, deep dives

3.4 Documentation requirements

- Decision log: every approval, rejection, or policy change recorded.
- Incident log: every incident, near-miss, or notable issue.
- Inventory: every AI system in use, with owner, purpose, data classes.
- Policy document: current and dated. Previous versions archived.
- Annual report: external-facing summary for stakeholders.

3.5 Authority and escalation

Make explicit who can:

- Approve a new AI tool (and at what cost threshold).
- Pause or terminate a deployed AI system.
- Override a policy in an emergency.
- Make public statements about AI use.
- Engage with regulators on AI matters.

4. Risk and Bias Assessment Toolkit

4.1 The risk assessment workflow

- 1 **Define the system.** What does it do? Who uses it? Who is affected?
- 2 **Identify risks.** Use the categories below. Add context-specific risks.
- 3 **Assess likelihood and impact.** Use a simple 1-5 scale unless you have something better.
- 4 **Prioritize.** Focus on high-impact risks first.
- 5 **Plan mitigations.** For each significant risk, name a mitigation owner and timeline.
- 6 **Document.** Risk register entries become part of the system's documentation.
- 7 **Review.** At least annually; sooner if conditions change.

4.2 Risk categories

- **Operational:** system fails, output is wrong, costs spiral.
- **Compliance:** regulation violated.
- **Reputation:** public embarrassment or trust loss.
- **Bias / fairness:** systematic harm to subgroups.
- **Privacy:** personal data exposed or misused.
- **Security:** system compromised.
- **Strategic:** wrong AI investment crowds out better options.
- **Workforce:** staff resistance, skills gap, role anxiety.

4.3 Bias assessment specifics

Treat bias as a discipline of its own. Common methods:

- **Demographic parity testing:** are outcomes similar across protected groups?
- **Equal opportunity testing:** are error rates similar across groups?
- **Counterfactual testing:** if you change a protected attribute, does the prediction change?
- **Stress testing:** how does the model behave on under-represented inputs?
- **Adversarial testing:** can prompts be crafted to elicit biased outputs?

4.4 Tools and techniques

- Documentation: model cards, system cards, data sheets.
- Quantitative tools: fairness metrics libraries, statistical disparate-impact tests.
- Qualitative tools: user research with affected communities.
- Process tools: red-teaming, bug bounties for bias.

5. Stakeholder Engagement Playbook

Stakeholder engagement is not the same as marketing. It is the practice of identifying everyone affected by an AI system and creating real channels for their voice.

5.1 Identifying stakeholders

- **Primary users:** people who interact with the AI directly.
- **Subjects:** people the AI makes decisions about.
- **Operators:** staff who run or maintain the AI.
- **Adjacent staff:** teams whose work is affected.
- **External:** regulators, partners, community groups, advocates.

5.2 Engagement levels (IAP2 spectrum)

Level	Goal	Methods
Inform	Provide information	Public notices, FAQs, town halls
Consult	Obtain feedback	Surveys, focus groups, public comment
Involve	Work directly with	Working groups, advisory boards
Collaborate	Partner in decisions	Co-design, joint working groups
Empower	Place decision in their hands	Citizen panels, governance roles

5.3 When to engage at each level

- **Low-stakes internal AI tools:** Inform staff.
- **Customer-facing AI:** Inform + Consult.
- **AI affecting employment or services:** Consult or Involve.
- **Public-sector AI affecting consequential decisions:** Involve or Collaborate.

5.4 Hard cases

- Engaging without creating false expectations.
- Handling vocal minorities versus silent majorities.
- When stakeholders disagree with each other.
- When stakeholders disagree with the data.

6. Implementation Roadmap Template

6.1 Phase 1 — Foundations (months 1-3)

- Adopt the six principles formally (board resolution).
- Designate accountability owner.
- Inventory current AI use.
- Draft policy. Adopt.
- Initial staff communication.

6.2 Phase 2 — Operating model (months 4-6)

- Stand up the AI Review Committee.
- Adopt the vendor approval workflow.
- Run first formal review of an active AI use.
- Conduct first risk and bias assessment on a deployed system.
- First quarterly board briefing.

6.3 Phase 3 — Maturation (months 7-12)

- Annual policy review.
- First annual public AI disclosure (if applicable).
- Stakeholder engagement on a significant deployment.
- Tabletop exercise: incident response drill.
- External review (consultant or peer review).

6.4 Phase 4 — Continuous improvement (year 2+)

- Update framework based on regulatory evolution.
- Refresh bias assessments as data shifts.
- Renew stakeholder relationships.
- Train new staff.
- Benchmark against peer organizations.

7. Quarterly Review Templates

7.1 Quarterly board briefing (one page)

Quarterly AI Program Review — Q[N] [Year]

Tools in production: [N] (list)

New approvals this quarter: [N]

Rejections / pauses: [N]

Incidents: [N] (severity)

Compliance status: [Green / Yellow / Red]

Risks emerging: [Brief description]

Decisions requested from board: [Specific items]

7.2 Annual program review

- Full inventory refresh.
- Policy review and update.
- Risk register review.
- Stakeholder engagement summary.
- External / peer review consideration.
- Strategic alignment check.
- Next year priorities.

7.3 Per-system annual review

- Original goals vs. actual outcomes.
- Bias and fairness re-assessment.
- Security re-assessment.
- Cost-benefit re-evaluation.
- Vendor relationship review.
- Decision: continue, modify, retire.

8. Maine-Specific Considerations

- **Demographics:** Maine's population is older and more rural than national averages. AI tools optimized for urban or younger users may underperform.
- **Languages:** French Acadian heritage and growing immigrant populations create language-coverage considerations.
- **Indigenous communities:** Wabanaki Nations have specific data sovereignty considerations.
- **Industries:** Strong nonprofit and municipal sectors with specific governance norms.
- **Regulatory environment:** Track state legislative activity on AI quarterly.
- **Workforce:** Maine's labor market has unique characteristics affecting AI hiring assumptions.

9. Self-Assessment Scorecard

Score yourself on each item: 0 (not started), 1 (in progress), 2 (in place but informal), 3 (formal and mature).

Area	Practice	Score
Direction	Six principles formally adopted	
Direction	Risk appetite stated	
Direction	Strategic AI plan	
Operating model	Accountability owner named	
Operating model	Review committee meeting regularly	
Operating model	Approved tool catalog maintained	
Operating model	Vendor approval workflow followed	
Risk & bias	Risk register maintained	
Risk & bias	Bias testing on consequential systems	
Risk & bias	Annual risk review	
Transparency	Policy published internally	
Transparency	External AI disclosure (if applicable)	
Transparency	Plain-language explanations available	
Engagement	Stakeholder mapping complete	
Engagement	Engagement plan for major systems	
Engagement	Feedback channels active	
Privacy & security	Data classification scheme	
Privacy & security	AI-specific security controls	
Privacy & security	Incident response plan	
Accountability	Decision log maintained	
Accountability	Annual board review	
Accountability	Consequences for violations defined	
	TOTAL (out of 66)	

Score interpretation

50-66: mature responsible AI program. 30-49: solid foundation, gaps to close. 15-29: early stage, prioritize quick wins. Below 15: start with section 6, Phase 1.

About AI Impact Maine

AI Impact Maine helps Maine businesses, municipalities, and nonprofits assess, plan, and responsibly implement AI. We offer readiness assessments, training programs, policy and governance design, vendor evaluation, and ongoing strategic support.

Get in touch

Email: ahmad@aiimpactmaine.com

Web: aiimpactmaine.com

Book a free 30-minute consultation: aiimpactmaine.com/contact